

mp1-jseytre

February 9, 2018

1 Simple classification

First we import the necessary packages and generate our datasets

```
In [69]: from mp1 import *
         from keras.models import Sequential, clone_model
         from keras.layers import Dense, Activation, Conv2D, Flatten, Dropout, MaxPooling2D
         from keras.utils import to_categorical
         from keras.regularizers import l2, l1
```

```
In [4]: [X_train1, Y_train1] = generate_dataset_classification(300, 20)
        Y_train1 = to_categorical(Y_train1)
```

Creating data:

```
0
100
200
```

1.0.1 First neural network (simple)

We then load our neural network: simple linear classifier with 3 neurons, using softmax for output, cross-entropy as loss and printing accuracy as gradient descent occurs.

```
In [5]: model1 = Sequential()
        model1.add(Dense(3, input_shape=(5184,), activation='softmax'))
        model1.compile(optimizer='sgd', loss='categorical_crossentropy',
                       metrics=['accuracy'])
```

```
In [6]: model1.fit(X_train1[0:200], Y_train1[0:200],
                   validation_data=(X_train1[200:300], Y_train1[200:300]),
                   epochs=20, batch_size=32)
```

Train on 200 samples, validate on 100 samples

Epoch 1/20

200/200 [=====] - 0s 2ms/step - loss: 8.4545 - acc: 0.2350 - val_loss

Epoch 2/20

200/200 [=====] - 0s 98us/step - loss: 10.2350 - acc: 0.3650 - val_loss

```

Epoch 3/20
200/200 [=====] - 0s 100us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 4/20
200/200 [=====] - 0s 95us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 5/20
200/200 [=====] - 0s 90us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 6/20
200/200 [=====] - 0s 88us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 7/20
200/200 [=====] - 0s 100us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 8/20
200/200 [=====] - 0s 93us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 9/20
200/200 [=====] - 0s 105us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 10/20
200/200 [=====] - 0s 90us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 11/20
200/200 [=====] - 0s 95us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 12/20
200/200 [=====] - 0s 95us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 13/20
200/200 [=====] - 0s 85us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 14/20
200/200 [=====] - 0s 93us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 15/20
200/200 [=====] - 0s 90us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 16/20
200/200 [=====] - 0s 95us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 17/20
200/200 [=====] - 0s 98us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 18/20
200/200 [=====] - 0s 90us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 19/20
200/200 [=====] - 0s 95us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350
Epoch 20/20
200/200 [=====] - 0s 100us/step - loss: 10.2350 - acc: 0.3650 - val_loss: 10.2350

```

```
Out [6]: <keras.callbacks.History at 0x1a461c49c18>
```

Instead of having to optimize the different settings of the Stochastic Gradient Descent, switching to Adam ensures good convergence.

```

In [7]: model1 = Sequential()
        model1.add(Dense(3, input_shape=(5184,), activation='softmax'))
        model1.compile(optimizer='adam', loss='categorical_crossentropy',
                       metrics=['accuracy'])

```

```
In [8]: model1.fit(X_train1[0:200], Y_train1[0:200],
```

```
validation_data=(X_train1[200:300], Y_train1[200:300]),  
epochs=20, batch_size=32)
```

Train on 200 samples, validate on 100 samples

Epoch 1/20

200/200 [=====] - 1s 3ms/step - loss: 2.0602 - acc: 0.3900 - val_loss:

Epoch 2/20

200/200 [=====] - 0s 103us/step - loss: 1.5625 - acc: 0.5400 - val_loss:

Epoch 3/20

200/200 [=====] - 0s 103us/step - loss: 1.1789 - acc: 0.6150 - val_loss:

Epoch 4/20

200/200 [=====] - 0s 98us/step - loss: 0.7508 - acc: 0.6800 - val_loss:

Epoch 5/20

200/200 [=====] - 0s 100us/step - loss: 0.5055 - acc: 0.7250 - val_loss:

Epoch 6/20

200/200 [=====] - 0s 98us/step - loss: 0.4010 - acc: 0.8900 - val_loss:

Epoch 7/20

200/200 [=====] - 0s 95us/step - loss: 0.3981 - acc: 0.8400 - val_loss:

Epoch 8/20

200/200 [=====] - 0s 95us/step - loss: 0.3292 - acc: 0.8900 - val_loss:

Epoch 9/20

200/200 [=====] - 0s 98us/step - loss: 0.2843 - acc: 0.9300 - val_loss:

Epoch 10/20

200/200 [=====] - 0s 125us/step - loss: 0.3516 - acc: 0.8500 - val_loss:

Epoch 11/20

200/200 [=====] - 0s 105us/step - loss: 0.3606 - acc: 0.8500 - val_loss:

Epoch 12/20

200/200 [=====] - 0s 110us/step - loss: 0.4067 - acc: 0.8050 - val_loss:

Epoch 13/20

200/200 [=====] - 0s 115us/step - loss: 0.2865 - acc: 0.8700 - val_loss:

Epoch 14/20

200/200 [=====] - 0s 105us/step - loss: 0.2134 - acc: 0.9200 - val_loss:

Epoch 15/20

200/200 [=====] - 0s 118us/step - loss: 0.1787 - acc: 0.9850 - val_loss:

Epoch 16/20

200/200 [=====] - 0s 110us/step - loss: 0.1621 - acc: 0.9950 - val_loss:

Epoch 17/20

200/200 [=====] - 0s 108us/step - loss: 0.1543 - acc: 0.9900 - val_loss:

Epoch 18/20

200/200 [=====] - 0s 108us/step - loss: 0.1503 - acc: 0.9750 - val_loss:

Epoch 19/20

200/200 [=====] - 0s 103us/step - loss: 0.1514 - acc: 0.9850 - val_loss:

Epoch 20/20

200/200 [=====] - 0s 105us/step - loss: 0.1272 - acc: 0.9950 - val_loss:

Out[8]: <keras.callbacks.History at 0x1a462218ba8>

1.0.2 Checking the classifier

We can check our classifier for all 3 classes

```
In [9]: X_test_r = generate_a_rectangle()
        X_test_d = generate_a_disk()
        X_test_t = generate_a_triangle()[0]
        X_test_r = X_test_r.reshape(1, X_test_r.shape[0])
        X_test_d = X_test_d.reshape(1, X_test_d.shape[0])
        X_test_t = X_test_t.reshape(1, X_test_t.shape[0])
        print(model1.predict(X_test_r))
        print(model1.predict(X_test_d))
        print(model1.predict(X_test_t))

[[1. 0. 0.]]
[[0. 1. 0.]]
[[0. 0. 1.]]
```

1.0.3 We then extract the weights and visualize them

We can even recognize the different shapes.

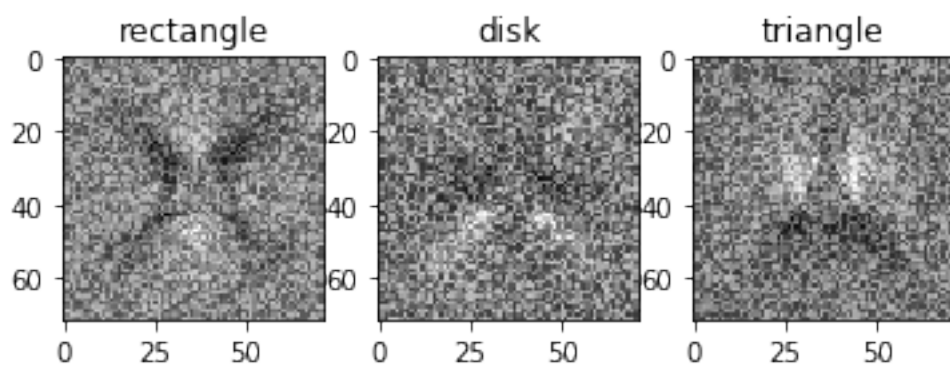
```
In [10]: weights = model1.get_weights()[0].reshape(72,72,3)

        plt.subplot(131)
        plt.imshow(weights[:, :, 0], cmap='gray')
        plt.title('rectangle')

        plt.subplot(132)
        plt.imshow(weights[:, :, 1], cmap='gray')
        plt.title('disk')

        plt.subplot(133)
        plt.imshow(weights[:, :, 2], cmap='gray')
        plt.title('triangle')

        plt.show()
```



2 A more complicated classification problem

First we generate the data and reshape it for the needs of the convolutional neural network (CNN). In order to have better results we use training and testing datasets that contain 1,000 images. Our CNN reaches a performance of: - 100% training accuracy - 93% testing at testing accuracy - 0.3 loss*

Note: implementing early stopping would have lowered the final *loss* at testing reached but seemingly not improved *accuracy*.

```
In [11]: n_samples = 1000
```

```
[X_train2, Y_train2] = generate_dataset_classification(n_samples, 20, True)
[X_test2, Y_test2] = generate_test_set_classification(n_samples, True)
Y_train2 = to_categorical(Y_train2)
X_train2 = X_train2.reshape(n_samples,72,72,1)
X_test2 = X_test2.reshape(n_samples,72,72,1)
```

Creating data:

0
100
200
300
400
500
600
700
800
900

Creating data:

0
100
200
300
400
500
600
700
800
900

2.0.1 Trying out the previous model for this more complicated task

Spoilers: it doesn't work (because of changing shapes and positions)

```
In [12]: X_train_temp = X_train2.reshape(n_samples, 5184)
        X_test_temp = X_test2.reshape(n_samples, 5184)
```

```
        model1.fit(X_train_temp, Y_train2,
                   validation_data=(X_test_temp, Y_test2),
                   epochs=20, batch_size=32)
```

Train on 1000 samples, validate on 1000 samples

```
Epoch 1/20
1000/1000 [=====] - 0s 108us/step - loss: 1.8027 - acc: 0.3640 - val_loss: 1.8027
Epoch 2/20
1000/1000 [=====] - 0s 98us/step - loss: 1.4284 - acc: 0.4370 - val_loss: 1.4284
Epoch 3/20
1000/1000 [=====] - 0s 95us/step - loss: 1.1621 - acc: 0.5090 - val_loss: 1.1621
Epoch 4/20
1000/1000 [=====] - 0s 98us/step - loss: 1.0360 - acc: 0.5170 - val_loss: 1.0360
Epoch 5/20
1000/1000 [=====] - 0s 98us/step - loss: 1.0424 - acc: 0.5070 - val_loss: 1.0424
Epoch 6/20
1000/1000 [=====] - 0s 99us/step - loss: 0.9215 - acc: 0.5760 - val_loss: 0.9215
Epoch 7/20
1000/1000 [=====] - 0s 99us/step - loss: 0.8243 - acc: 0.6050 - val_loss: 0.8243
Epoch 8/20
1000/1000 [=====] - 0s 100us/step - loss: 0.8862 - acc: 0.5750 - val_loss: 0.8862
Epoch 9/20
1000/1000 [=====] - 0s 102us/step - loss: 0.9647 - acc: 0.5460 - val_loss: 0.9647
Epoch 10/20
1000/1000 [=====] - 0s 99us/step - loss: 0.8568 - acc: 0.5870 - val_loss: 0.8568
Epoch 11/20
1000/1000 [=====] - 0s 100us/step - loss: 0.9198 - acc: 0.5780 - val_loss: 0.9198
Epoch 12/20
1000/1000 [=====] - 0s 108us/step - loss: 0.7719 - acc: 0.6460 - val_loss: 0.7719
Epoch 13/20
1000/1000 [=====] - 0s 104us/step - loss: 0.7770 - acc: 0.6090 - val_loss: 0.7770
Epoch 14/20
1000/1000 [=====] - 0s 95us/step - loss: 1.0643 - acc: 0.5840 - val_loss: 1.0643
Epoch 15/20
1000/1000 [=====] - 0s 99us/step - loss: 0.9323 - acc: 0.5770 - val_loss: 0.9323
Epoch 16/20
1000/1000 [=====] - 0s 98us/step - loss: 0.7773 - acc: 0.6200 - val_loss: 0.7773
Epoch 17/20
1000/1000 [=====] - 0s 97us/step - loss: 0.7959 - acc: 0.6110 - val_loss: 0.7959
Epoch 18/20
1000/1000 [=====] - 0s 100us/step - loss: 0.8007 - acc: 0.6280 - val_loss: 0.8007
Epoch 19/20
1000/1000 [=====] - 0s 103us/step - loss: 0.9113 - acc: 0.6070 - val_loss: 0.9113
Epoch 20/20
1000/1000 [=====] - 0s 114us/step - loss: 0.9285 - acc: 0.6110 - val_loss: 0.9285
```

```
Out [12]: <keras.callbacks.History at 0x1a46ae25c18>
```

2.0.2 The final CNN used

64 filters of size 3x3 with 4x4 MaxPooling, 0.5 Dropout and a fully-connected 124-neuron hidden layer. Loss: cross-entropy; optimizer: adam.

```
In [13]: model2 = Sequential()
          model2.add(Conv2D(filters=64, kernel_size=(3,3), input_shape=(72,72,1), activation='relu'))
          model2.add(MaxPooling2D((4,4)))
          model2.add(Dropout(0.5))
          model2.add(Flatten())
          model2.add(Dense(124, activation='relu'))
          model2.add(Dense(3, activation='softmax'))
          model2.compile(optimizer='adam', loss='categorical_crossentropy',
                        metrics=['accuracy'])
          print(model2.summary())
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 70, 70, 64)	640
max_pooling2d_1 (MaxPooling2D)	(None, 17, 17, 64)	0
dropout_1 (Dropout)	(None, 17, 17, 64)	0
flatten_1 (Flatten)	(None, 18496)	0
dense_3 (Dense)	(None, 124)	2293628
dense_4 (Dense)	(None, 3)	375

Total params: 2,294,643
Trainable params: 2,294,643
Non-trainable params: 0

```
In [14]: model2.fit(X_train2, Y_train2,
                    validation_data=(X_test2, Y_test2),
                    epochs=50, batch_size=32)
```

Train on 1000 samples, validate on 1000 samples

Epoch 1/50

1000/1000 [=====] - 9s 9ms/step - loss: 1.0912 - acc: 0.5030 - val_loss: 1.0912 - acc: 0.5030

Epoch 2/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.7679 - acc: 0.6990 - val_loss: 0.7679
Epoch 3/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.6664 - acc: 0.7610 - val_loss: 0.6664
Epoch 4/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.5392 - acc: 0.8240 - val_loss: 0.5392
Epoch 5/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.4681 - acc: 0.8420 - val_loss: 0.4681
Epoch 6/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.4346 - acc: 0.8620 - val_loss: 0.4346
Epoch 7/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.3590 - acc: 0.8850 - val_loss: 0.3590
Epoch 8/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.2953 - acc: 0.9050 - val_loss: 0.2953
Epoch 9/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.2467 - acc: 0.9350 - val_loss: 0.2467
Epoch 10/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.2190 - acc: 0.9350 - val_loss: 0.2190
Epoch 11/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.1928 - acc: 0.9440 - val_loss: 0.1928
Epoch 12/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.1652 - acc: 0.9580 - val_loss: 0.1652
Epoch 13/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.1579 - acc: 0.9540 - val_loss: 0.1579
Epoch 14/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.1347 - acc: 0.9620 - val_loss: 0.1347
Epoch 15/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.1043 - acc: 0.9720 - val_loss: 0.1043
Epoch 16/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.1090 - acc: 0.9700 - val_loss: 0.1090
Epoch 17/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.1158 - acc: 0.9630 - val_loss: 0.1158
Epoch 18/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0735 - acc: 0.9850 - val_loss: 0.0735
Epoch 19/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0731 - acc: 0.9820 - val_loss: 0.0731
Epoch 20/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0599 - acc: 0.9910 - val_loss: 0.0599
Epoch 21/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0562 - acc: 0.9860 - val_loss: 0.0562
Epoch 22/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0517 - acc: 0.9910 - val_loss: 0.0517
Epoch 23/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0563 - acc: 0.9850 - val_loss: 0.0563
Epoch 24/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0389 - acc: 0.9960 - val_loss: 0.0389
Epoch 25/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0388 - acc: 0.9900 - val_loss: 0.0388

Epoch 26/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0310 - acc: 0.9970 - val_loss: 0.0310
Epoch 27/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0284 - acc: 0.9980 - val_loss: 0.0284
Epoch 28/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0310 - acc: 0.9980 - val_loss: 0.0310
Epoch 29/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0326 - acc: 0.9900 - val_loss: 0.0326
Epoch 30/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0314 - acc: 0.9950 - val_loss: 0.0314
Epoch 31/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0313 - acc: 0.9950 - val_loss: 0.0313
Epoch 32/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0263 - acc: 0.9940 - val_loss: 0.0263
Epoch 33/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0213 - acc: 0.9960 - val_loss: 0.0213
Epoch 34/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0165 - acc: 0.9980 - val_loss: 0.0165
Epoch 35/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0167 - acc: 0.9990 - val_loss: 0.0167
Epoch 36/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0177 - acc: 0.9990 - val_loss: 0.0177
Epoch 37/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0145 - acc: 0.9980 - val_loss: 0.0145
Epoch 38/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0217 - acc: 0.9960 - val_loss: 0.0217
Epoch 39/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0194 - acc: 0.9960 - val_loss: 0.0194
Epoch 40/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0150 - acc: 0.9990 - val_loss: 0.0150
Epoch 41/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0120 - acc: 1.0000 - val_loss: 0.0120
Epoch 42/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0094 - acc: 0.9990 - val_loss: 0.0094
Epoch 43/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0111 - acc: 0.9990 - val_loss: 0.0111
Epoch 44/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0098 - acc: 0.9990 - val_loss: 0.0098
Epoch 45/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0080 - acc: 1.0000 - val_loss: 0.0080
Epoch 46/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0103 - acc: 1.0000 - val_loss: 0.0103
Epoch 47/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0086 - acc: 0.9990 - val_loss: 0.0086
Epoch 48/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0098 - acc: 0.9990 - val_loss: 0.0098
Epoch 49/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0093 - acc: 0.9980 - val_loss: 0.0093

```
Epoch 50/50
1000/1000 [=====] - 6s 6ms/step - loss: 0.0064 - acc: 1.0000 - val_loss:
```

```
Out[14]: <keras.callbacks.History at 0x1a46ae25278>
```

```
In [15]: model2.evaluate(X_test2, Y_test2)
```

```
1000/1000 [=====] - 1s 994us/step
```

```
Out[15]: [0.27490175648313014, 0.925]
```

3 Regression problem

For this regression problem, we had to sort the vertices dataset in a way that we are going to describe. In the end we noticed that the task the CNN is trying to accomplish can necessitate different filters based on what vertex it wants to locate. To that end we ended up training a different CNN (that remains pretty simple in terms of architecture) for each of the 6 coordinates.

Note: we could have tried training only 3 different CNNs and seen if the same CNN can serve for x and y coordinates.

First we load the dataset and visualize an example.

```
In [16]: [X_train3, Y_train3] = generate_dataset_regression(300, 20)
[X_test3, Y_test3] = generate_test_set_regression()
X_train3 = X_train3.reshape(300,72,72,1)
X_test3 = X_test3.reshape(300,72,72,1)

i=2
visualize_prediction(X_train3[i], Y_train3[i])
```

```
Creating data:
```

```
0
```

```
100
```


```
200
```

```
Creating data:
```

```
0
```

```
100
```

```
200
```



3.0.1 Normalizing the dataset

We can notice that the vertices are not sorted in the dataset at hand. Here we implement a function that sorts the vertices by computing the projection of the vertices on the (1,1) vector. We also subtract 0.5 to the values so that the values are centered on zero (better initialization for the network).

Additionally, the function returns only one of the coordinates so that we can treat them all separately.

```
In [17]: def sort_vertices(dataset, index):
    middle = 0.5
    # res = np.zeros((dataset.shape))
    res = np.zeros((300, 1))
    for i in range(0, 300):
        temp1 = [dataset[i,2*m]+dataset[i,2*m+1] for m in range(0, 3)]
        j = np.argmax(temp1)
        k = np.argmin(temp1)
        l = 3 - j - k
        temp2 = [dataset[i,2*j]-middle, dataset[i,2*j+1]-middle,
                  dataset[i,2*k]-middle, dataset[i,2*k+1]-middle,
                  dataset[i,2*l]-middle, dataset[i,2*l+1]-middle]
```

```

        res[i,:] = [temp2[index]]
    return res

```

```

In [18]: Y_train4 = [sort_vertices(Y_train3, i) for i in range(0, 6)]
        Y_test4 = [sort_vertices(Y_test3, i) for i in range(0, 6)]

```

3.0.2 CNN architecture

The model used is the same for all coordinates: 32 filters of size 9x9, on top of which we add a 32-neuron fully-connected layer. Activation: ReLu; loss: mean squared error; optimizer: Adam.

```

In [129]: reg_param = 0.00001

```

```

model3 = Sequential()
model3.add(Conv2D(filters=32, kernel_size=(5,5), input_shape=(72,72,1), activation='relu'))
model3.add(Conv2D(filters=32, kernel_size=(3,3), activation='relu', kernel_regularizer=l2(reg_param)))
# model3.add(Conv2D(filters=32, kernel_size=(2,2), activation='relu', kernel_regularizer=l2(reg_param)))
model3.add(Flatten())
model3.add(Dense(64, activation='relu', kernel_regularizer=l2(reg_param)))
model3.add(Dense(1))
# model3.add(Dense(6))
model3.compile(optimizer='adam', loss='mean_squared_error')
print(model3.summary())

```

```

-----
Layer (type)                 Output Shape              Param #
-----
conv2d_82 (Conv2D)           (None, 68, 68, 32)       832
-----
conv2d_83 (Conv2D)           (None, 66, 66, 32)       9248
-----
flatten_40 (Flatten)         (None, 139392)           0
-----
dense_77 (Dense)              (None, 64)                8921152
-----
dense_78 (Dense)              (None, 1)                  65
=====
Total params: 8,931,297
Trainable params: 8,931,297
Non-trainable params: 0
-----
None

```

3.0.3 Training our 6 CNNs

Here we train our 6 CNNs one after the other. Given the highly non-convex loss function, several runs are necessary to obtain a good performance for each coordinate.

For that reason we retrain the models whose performance is not deemed good enough (criteria: validation loss > 0.05).

```

In [135]: threshold = 0.05
retrain_from_scratch = False
errors = np.ones((6,2))
for i in range(0, 6):
    try:
        train_error = models[i].evaluate(X_train3, Y_train4[i], verbose = 0)
        test_error = models[i].evaluate(X_test3, Y_test4[i], verbose = 0)
        print("Evaluating model %i: train error %.3f - test error %.3f" % (i, train_error, test_error))
    except:
        try:
            print("\nModel %i/%i was not defined..." % (i, len(models)))
        except:
            print("\nModels not defined!")
            models = [clone_model(model3) for _ in range(0, 6)]
            test_error = 1
if test_error > threshold:
    print("\n\nRetraining model %i because it was above threshold (%.3f vs %.3f)" % (i, train_error, test_error))
    if retrain_from_scratch:
        models[i] = clone_model(model3)
        models[i].compile(optimizer='adam', loss='mean_squared_error')
        models[i].fit(X_train3, Y_train4[i], validation_data=(X_test3, Y_test4[i]), epochs=20, batch_size=64)
    print("\n\n")
    train_error = models[i].evaluate(X_train3, Y_train4[i], verbose = 0)
    test_error = models[i].evaluate(X_test3, Y_test4[i], verbose = 0)
    print("Evaluating model %i: train error %.3f - test error %.3f" % (i, train_error, test_error))
errors[i,:] = [train_error, test_error]
print("\nDone evaluating models!")

```

```

Evaluating model 0: train error 0.009 - test error 0.030
Evaluating model 1: train error 0.007 - test error 0.028
Evaluating model 2: train error 0.014 - test error 0.037
Evaluating model 3: train error 0.012 - test error 0.036
Evaluating model 4: train error 0.003 - test error 0.047
Evaluating model 5: train error 0.003 - test error 0.050

```

Done evaluating models!

```

In [53]: model3.fit(X_train3, Y_train4,
                    validation_data=(X_test3, Y_test4),
                    epochs=20, batch_size=64)

```

Train on 300 samples, validate on 300 samples

Epoch 1/20

300/300 [=====] - 4s 13ms/step - loss: 27.3224 - val_loss: 0.0837

Epoch 2/20

300/300 [=====] - 3s 9ms/step - loss: 0.0811 - val_loss: 0.0838

```

Epoch 3/20
300/300 [=====] - 3s 9ms/step - loss: 0.0812 - val_loss: 0.0838
Epoch 4/20
300/300 [=====] - 3s 9ms/step - loss: 0.1034 - val_loss: 0.0837
Epoch 5/20
300/300 [=====] - 3s 9ms/step - loss: 0.0810 - val_loss: 0.0836
Epoch 6/20
300/300 [=====] - 3s 10ms/step - loss: 0.0809 - val_loss: 0.0834
Epoch 7/20
300/300 [=====] - 3s 9ms/step - loss: 0.0807 - val_loss: 0.0832
Epoch 8/20
300/300 [=====] - 3s 9ms/step - loss: 0.0805 - val_loss: 0.0830
Epoch 9/20
300/300 [=====] - 3s 10ms/step - loss: 0.0802 - val_loss: 0.0828
Epoch 10/20
300/300 [=====] - 3s 10ms/step - loss: 0.0800 - val_loss: 0.0825
Epoch 11/20
300/300 [=====] - 3s 9ms/step - loss: 0.0797 - val_loss: 0.0823
Epoch 12/20
300/300 [=====] - 3s 9ms/step - loss: 0.0795 - val_loss: 0.0820
Epoch 13/20
300/300 [=====] - 3s 10ms/step - loss: 0.0792 - val_loss: 0.0818
Epoch 14/20
300/300 [=====] - 3s 9ms/step - loss: 0.0789 - val_loss: 0.0815
Epoch 15/20
300/300 [=====] - 4s 12ms/step - loss: 0.0786 - val_loss: 0.0812
Epoch 16/20
300/300 [=====] - 3s 10ms/step - loss: 0.0783 - val_loss: 0.0809
Epoch 17/20
300/300 [=====] - 3s 9ms/step - loss: 0.0780 - val_loss: 0.0807
Epoch 18/20
300/300 [=====] - 3s 9ms/step - loss: 0.0777 - val_loss: 0.0804
Epoch 19/20
300/300 [=====] - 3s 10ms/step - loss: 0.0774 - val_loss: 0.0801
Epoch 20/20
300/300 [=====] - 3s 10ms/step - loss: 0.0771 - val_loss: 0.0798

```

Out [53]: <keras.callbacks.History at 0x281daa87518>


3.0.4 Visualization of the results


The results are decent. There is quite a bit of overfitting: I experimented with L1, L2 regularization as well as dropout. With this setup I manage to get under 0.05 validation loss on all coordinates.


I believe they could be improved by maybe figuring out a better architecture (how can we regularize?) and by expanding the dataset.


```
In [136]: for i in range(0,5):
```


```
predictions = np.array([models[k].predict(X_train3)[i,0]+0.5 for k in range(0,6)])  
visualize_prediction(X_train3[i,:,:], predictions)
```











```
In [137]: for i in range(0,5):  
           predictions = np.array([models[k].predict(X_test3)[i,0]+0.5 for k in range(0,6)])  
           visualize_prediction(X_test3[i,:,:], predictions)
```

